

submissions

Q2

```
https://github.com/Aravindkumar-Rajendran/ERA-v3-session-9/blob/main/training-g5.12xlarge.log
```

```
https://github.com/Aravindkumar-Rajendran/ERA-v3-session-9/blob/main/training-g4dn.2xlarge.log
```

We have started training the models with g4dn.2xlarge and that took around 40 hours for completing 10 epochs (close to 3.5 hours for 1 epoch). We have to choose this because we only had 16 vCPUs in our service quota.

I applied for more service quota while starting the training on above instance but did not receive it on time.

Then I found that I had some more quota in other region. I moved the EBS volume to that region and started an instance with g5.12xlarge with 4 GPUs. I modified the code with multi-gpu support and started the training. I was taking around 30 minutes per epoch and ran for 15 epochs and achieved more than 40% accuracy. And we had to stop it there because it was already end of the day of 31st Dec (we only had credits till the end of the year).

Q3

Share the screenshot of your EC2 usage

Your Answer:

<https://github.com/Aravindkumar-Rajendran/ERA-v3-session-9/blob,>

<https://github.com/Aravindkumar-Rajendran/ERA-v3-session-9/blob,>

Q4

42.57

Q5

<https://huggingface.co/spaces/kishkath/Res-Imagenet>

Q6

<https://github.com/Aravindkumar-Rajendran/ERA-v3-session-9/tree,>